Samantha An

ACM Research Lab

October 27, 2026

The Mythos of Model Interpretability

In The Mythos of Model Interpretability, Zachary C. Lipton discusses the lack of a clear definition of interpretability in machine learning research. Although researchers often argue that interpretable models are preferable because they are easier to understand and trust, Lipton points out that interpretability can mean different things depending on the goal. Some researchers focus on understanding how a model works internally, while others focus on explaining its predictions after it has been trained. The paper argues that interpretability should not be treated as a single property that automatically makes a model better. One of the main ideas in the paper is the distinction between transparency and post-hoc interpretability. Transparency refers to understanding a model's internal structure and learning process. For example, a simple decision tree may be easier to follow than a deep neural network. Post-hoc interpretability, on the other hand, involves using methods such as visualizations, local explanations, or similar examples to understand a model's predictions without necessarily knowing exactly how it works internally.

However, Lipton warns that these explanations can be misleading. An explanation may appear reasonable to a human without accurately reflecting the model's actual decision-making process. Another important point is that predictive accuracy does not necessarily mean a model is trustworthy or suitable for real-world use. A model can perform well on test data while still producing unfair decisions, failing when its environment changes, or learning correlations that do not represent causal relationships. For example, a model might identify an association between a

medical condition and patient survival without recognizing that the relationship is influenced by differences in medical treatment. This shows why researchers need to consider what they actually want a model to accomplish beyond simply making accurate predictions.

Overall, I found the paper important because it challenges the assumption that simpler models are always more interpretable or that explanations automatically make complex models trustworthy. I think the most valuable takeaway is that researchers need to define what they mean by interpretability before claiming that a model achieves it. When developing or evaluating machine learning models, we should ask whether an explanation is faithful to the model, whether it provides useful information, and whether it addresses the specific problem we are trying to solve. Rather than treating interpretability as an advantage by itself, researchers should evaluate how well it serves the needs of the people using the model.